



Playing with time series data in python





Nele Custers
nele.custers@pxl.be

Year	Breakthroughs in AI	Datasets (First Available)	Algorithms (First Proposed)
1994	Human-level spontaneous speech recognition	Spoken Wall Street Journal articles and other texts (1991)	Hidden Markov Model (1984)
1997	IBM Deep Blue defeated Garry Kasparov	700,000 Grandmaster chess games, aka "The Extended Book" (1991)	Negascout planning algorithm (1983)
2005	Google's Arabic- and Chinese-to-English translation	1.8 trillion tokens from Google Web and News pages (collected in 2005)	Statistical machine translation algorithm (1988)
2011	IBM Watson became the world Jeopardy! champion	8.6 million documents from Wikipedia, Wiktionary, Wikiquote, and Project Gutenberg (updated in 2010)	Mixture-of-Experts algorithm (1991)
2014	Google's GoogLeNet object classification at near-human performance	ImageNet corpus of 1.5 million labeled images and 1,000 object categories (2010)	Convolution neural network algorithm (1989)
2015	Google's Deepmind achieved human parity in playing 29 Atari games by learning general control from video	Arcade Learning Environment dataset of over 50 Atari games (2013)	Q-learning algorithm (1992)
Average No. of Years to Breakthrough:		3 years	18 years

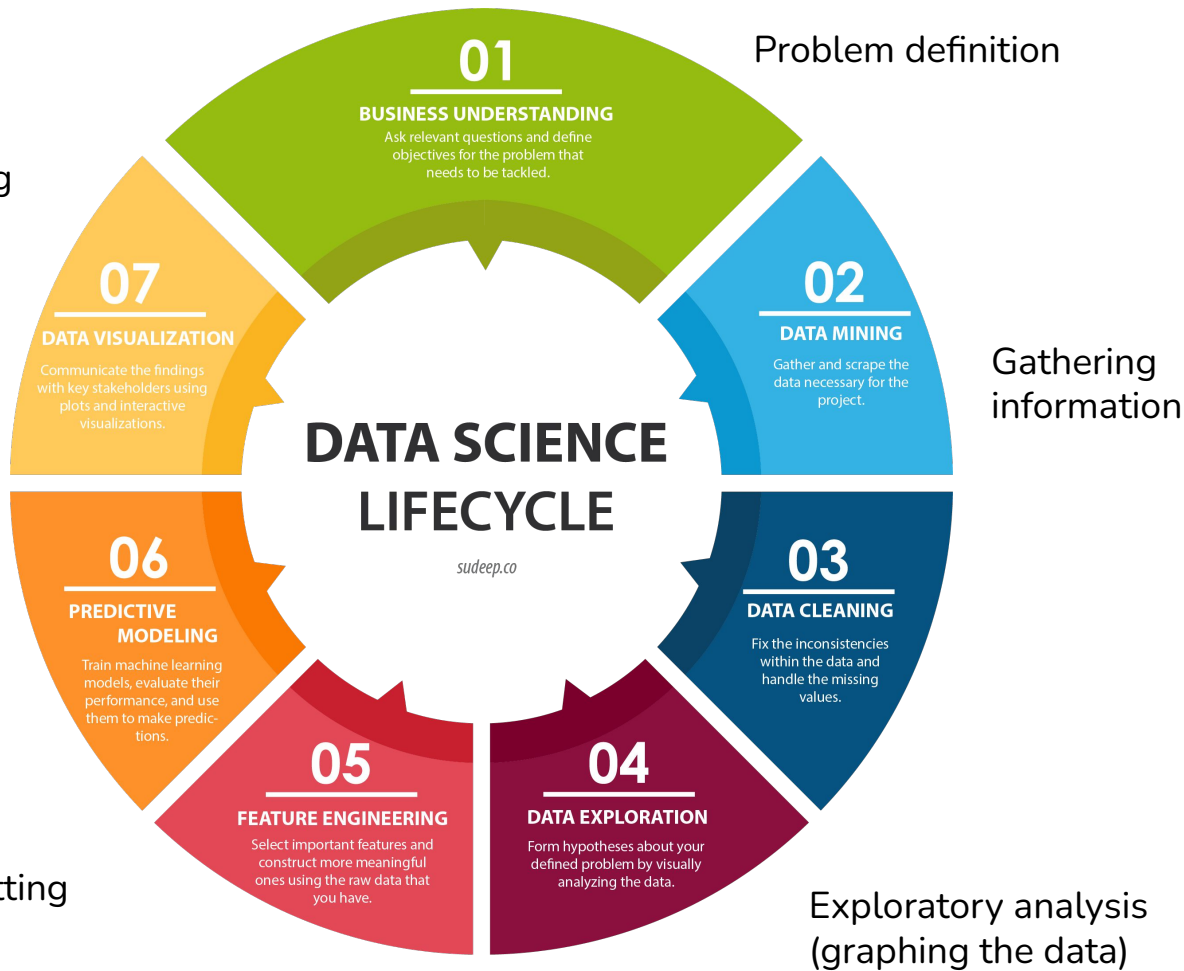


What is a Time Series?

***Time series** is a sequence of observations recorded at regular time intervals.*



Using and evaluating
a forecasting model





Time series patterns

Trend

A *trend* exists when there is a long-term increase or decrease in the data.

Seasonal


A *seasonal* pattern occurs when a time series is affected by seasonal factors such as the time of the year or the day of the week.

Cyclic


A *cycle* occurs when the data exhibit rises and falls that are not of a fixed frequency.

Time series patterns


Sales of new one-family houses, USA




US treasury bill contracts



Australian quarterly electricity production




Google daily changes in closing stock price






Exploratory analysis

*Perfect
Positive
Correlation*




1

*High
Positive
Correlation*




0.9

*Low
Positive
Correlation*



0.5

*No
Correlation*



0

*Low
Negative
Correlation*




-0.5

*High
Negative
Correlation*



-0.9

*Perfect
Negative
Correlation*



-1



Python libraries

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

pandas provides data structures for in-memory analytics (which makes **using pandas to analyze datasets that are larger than memory datasets somewhat tricky**).

matplotlib offers static, animated and interactive visualizations in Python.

seaborn is a Python data visualization library based on matplotlib.

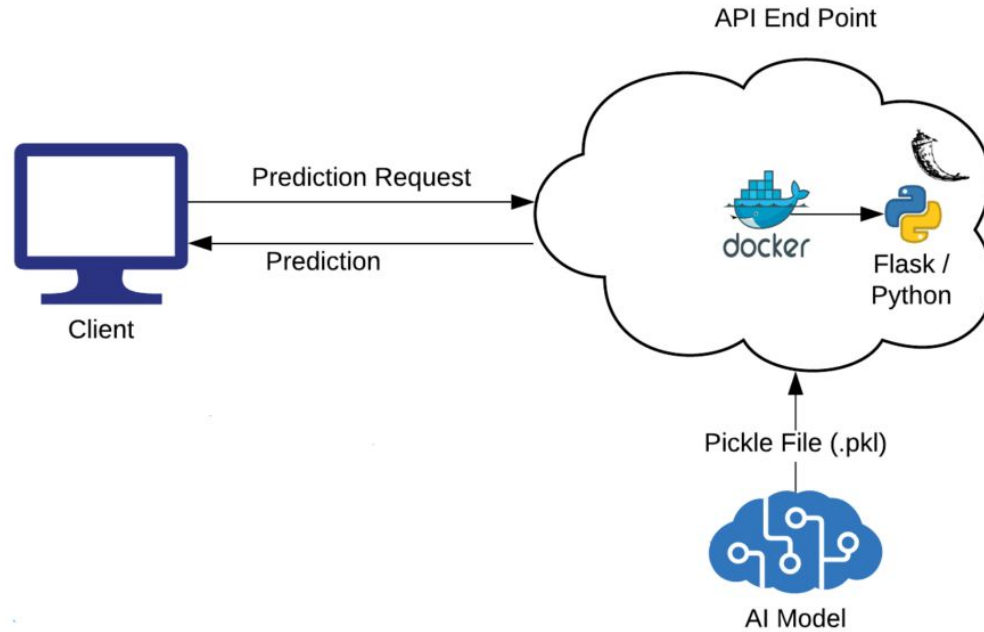
numpy is the standard for working with numerical data in Python.

Several solutions for forecasting



PROPHET

Deploying





Resources

Forecasting principles and practice: <https://otexts.com/fpp2/>

<https://facebook.github.io/prophet/>

<https://towardsdatascience.com/playing-with-time-series-data-in-python-959e2485bff8>

<https://towardsdatascience.com/time-series-analysis-with-theory-plots-and-code-part-1-dd3ea417d8c4>

<https://www.digitalocean.com/community/tutorials/how-to-build-and-deploy-a-flask-application-using-docker-on-ubuntu-20-04>

<https://blog.jetbrains.com/pycharm/2017/03/docker-compose-getting-flask-up-and-running/>

<http://sudeep.co/data-science/Understanding-the-Data-Science-Lifecycle/>

<https://www.vantage-ai.com/en/blog/4-strategies-how-to-deal-with-large-datasets-in-pandas>